

Selected Topics in Nature-inspired Algorithms

Seminar

Winter 2017

Nico Potyka

- Introduction
- Organization
- Optimization and Metaheuristics
- Local Search

Introduction




Nature-inspired Algorithms

- nature-inspired algorithms take inspiration from problem solving methods inherent to biological systems
- they are often easy to understand conceptually, but can be difficult to tune
- some classical approaches are
 - Genetic Algorithms
 - Differential Evolution
 - Ant Colony Optimization
 - Particle Swarm Optimization



Applications

- successful applications include
 - Discrete and Numerical Optimization (logistics, routing)
 - Constraint Satisfaction (frequency assignment, scheduling)
 - Machine Learning (intrusion detection, anomaly detection)



Development of the Field

- the field keeps on developing **new nature-inspired algorithms**

- including **successful ones**

- Firefly Algorithm
- Bee Colony Optimization
- Bat-inspired Algorithms



- and **rather exotic ones**

- Reincarnation Algorithm
- Zombie Survival Optimization
- Community of Scientists Opt.



A Critical View

- most algorithms come down to a **few principles** like
 - adaptation
 - natural selection
 - swarm behavior
- **lack of novelty** is often hidden beneath biological metaphors
- as put in the **Evolutionary Computation Bestiary**¹, the field is

[...] as The island of Doctor Moreau: a place with a few good creatures, but which are vastly outnumbered by mindless beasts.

¹<https://github.com/fcampelo/EC-Bestiary>

On the bright Side

- many nature-inspired algorithms have **proven useful for solving diverse problems** heuristically
- **no free lunch theorems**² justify diversity of the field
- intuitively, these theorem show that there can be no algorithm that outperforms all other algorithms for every problem



²Wolpert, D. H., & Macready, W. G. No free lunch theorems for optimization. IEEE transactions on evolutionary computation, 1(1), 67-82: 1997.

Goals of the Course

- our goal is to understand a selection of **well established nature-inspired algorithms** properly
- the algorithms are usually **simple to understand** conceptually
- however, they can usually be configured by a **significant number of parameters**
- tuning these parameters is often art rather than science
- we will therefore have **programming tasks** where you
 - 1 implement algorithms
 - 2 and tune the parameters for a given problem

Organization

Prerequisites

- basic mathematical concepts
- basic computer science concepts
- basic programming skills (can be developed during course)

Course Structure

- we will alternate theory and practice sessions
- **Theory Session:** introduction to one framework and subsequent discussion
- **Practice Session:** presentation and discussion of experiences with implementation and parameter tuning

Preparation for Theory Session (presenting group)

- familiarize yourself with topic
- find additional [citable literature](#) if necessary
(use Google Scholar for instance)
- prepare [well-structured, correct and reader-friendly slides](#)
- back up claims with [references to the literature](#)
- prepare a few [questions for discussion](#)
- presentation length based on group size

Preparation for Practice Session (every group)

- read slides/ basic literature for framework
- implement algorithm for given problem
- make experiments with different parameters
- report your findings in short presentation
- every group member has to present findings for at least one implementation
- every group member must be able to answer questions

Grading based on

① presentations

- reader-friendly slides
- self-contained, comprehensible presentation
- well prepared for questions and discussion

② participation in discussion

- raise critical questions
- argue objectively

③ programming exercises

- systematic parameter tuning
- clean and well documented code

(first few programming exercises will have less weight)

Topics and Basic Literature:

- 26/10: Introduction and Basic Local Search Algorithms

- 09/11: Genetic Algorithms

Sastry, K., Goldberg, D. E., & Kendall, G. Genetic algorithms. In Search methodologies (pp. 93-117). Springer US: 2014.

- 23/11: Ant Colony Optimization

Merkle, D., & Middendorf, M. Swarm intelligence. In Search methodologies (pp. 213-242). Springer US: 2014.

- 07/12: Differential Evolution

Das, S., & Suganthan, P. N. Differential evolution: A survey of the state-of-the-art. IEEE transactions on evolutionary computation, 15(1), 4-31: 2011.

- 21/12: Bee Colony Optimization

Teodorović, D. Bee colony optimization (BCO). Innovations in swarm intelligence, 39-60: 2009.

- 18/01: Particle Swarm Optimization

Poli, R., Kennedy, J., & Blackwell, T. Particle swarm optimization. Swarm intelligence, 1(1), 33-57: 2007.

- 01/02: Artificial Immune Systems

Aickelin, U., Dasgupta, D., & Gu, F. Artificial immune systems. In Search Methodologies (pp. 187-211). Springer US: 2014.

Groups and Topics:


- 6 topics: 6 or 12 groups
(dependent on number of participants)
- send me your preferred group members and topics until Monday morning
- if you have little programming experience, you can focus on parameter tuning initially and learn during the course

Optimization Problems and Meta-Heuristics

Some interesting Classes of Algorithms

- during the next weeks, we will in particular look at optimization problems
- in this context, nature-inspired algorithms can be regarded as meta-heuristics
- meta-heuristics can be seen as a special class of algorithms
 - **Optimal Algorithms:** solve problem optimally
 - **Approximation Algorithms:** solve problem with guaranteed error bound (e.g., at most 20% worse than optimal)
 - **Heuristics:** solve problem often well in practice, but there are no general guarantees (error can be arbitrarily large)
 - **Meta-Heuristics:** template heuristics that can be applied to many problems

Why Heuristics?



- many interesting problems **cannot be solved efficiently** under the usual complexity-theoretical assumptions
- in fact, many **cannot even be approximated efficiently**
- for many of these problems, **heuristics** have been developed that **'often' perform well**
- **meta-heuristics** abstract from a concrete problem and **can be applied to many problems**

Optimization Problems

- **optimization problems** have the following ingredients
 - variables with corresponding domains
 - constraints on values that variables can take (optional)
 - objective function (cost or utility of variable assignments)


$$\text{minimize} \quad \sum_{i \in F, j \in C} c_{ij} x_{ij} + \sum_{i \in F} f_i y_i$$

$$\begin{aligned} \text{subject to} \quad & \sum_{i \in F} x_{ij} \geq 1, & j \in C \\ & y_i - x_{ij} \geq 0, & i \in F, j \in C \\ & x_{ij} \in \{0, 1\}, & i \in F, j \in C \\ & y_i \in \{0, 1\}, & i \in F \end{aligned}$$

$$\begin{aligned} \min \quad & \sum_{i \in I} \left(\sum_{i \in I} f_i^t y_i^t + \sum_{i \in I} \sum_{j \in J} c_{ij}^t x_{ij} + \sum_{j \in J} h_j^t z_j^t \right), \\ \text{subject to} \quad & \sum_{i \in I} x_{ij}^t + z_j^t = d_j^t, \quad t \in T, j \in J, \\ & \sum_{j \in J} x_{ij}^t \leq q_i y_i^t, \quad t \in T, i \in I, \\ & \sum_{i \in I} y_i^t \leq p^t, \quad t \in T, \\ & y_i^t \leq y_i^{t+1}, \quad t \in T \setminus \{k\}, i \in I, \\ & y_i^t \in \{0, 1\}, \quad t \in T, i \in I, \\ & x_{ij}^t \geq 0, \quad t \in T, i \in I, j \in J \\ & z_j^t \geq 0, \quad t \in T, j \in J. \end{aligned}$$

- **examples** include
 - loss minimization in ML (unconstrained)
 - routing problems (discrete)
 - logistic problems (discrete or numerical)

Traveling Salesman Problem




- **Problem:** given complete weighted graph, compute cheapest cycle that contains all nodes
- **Examples**
 - find the shortest tour visiting all offices in the US that starts and ends in your own office (minimize **travel cost**)
 - find the shortest tour visiting all tourist attractions starting and ending in your hotel (minimize **walking distance**)
 - find the shortest tour visiting all customers starting and ending in your distribution center (minimize **time**)

Traveling Salesman Problem Formalized

- given **completely connected weighted graph** (V, E, c) with
 - n nodes $V = \{1, \dots, n\}$
 - edges $E = V \times V$
 - cost function $c : E \rightarrow \mathbb{R}_0^+$
- compute shortest cycle starting from 1 and visiting all nodes
- **optimization problem**
 - Variables: $n - 1$ variables with domain $\{2, \dots, n\}$
($x_i = j$ means that we visit node j at time i)
 - no constraints
 - Objective function: cost of tour
$$f(x_1, \dots, x_{n-1}) = c(1, x_1) + c(x_{n-1}, 1) + \sum_{i=1}^{n-2} c(x_i, x_{i+1})$$

Example

c(row,col)	1	2	3	4	5
1	0	2	3	4	5
2	2	0	3	2	3
3	3	3	0	1	1
4	4	2	1	0	2
5	5	3	1	2	0



We abbreviate assignment $(x_1 = v_1, \dots, x_5 = v_5)$ by (v_1, \dots, v_5)

- $f(2, 3, 4, 5) = 2 + 3 + 1 + 2 + 5 = 13$
- $f(3, 4, 5, 2) = 3 + 1 + 2 + 3 + 2 = 11$
- $f(5, 2, 3, 4) = 5 + 3 + 3 + 1 + 4 = 16$

Complexity of TSP

- TSP is NP-complete
- we cannot even find an efficient approximation algorithm if $P \neq NP$, see (Vazirani 2013)³, Theorem 3.6 for a proof
- if the edge costs satisfy triangle inequality ([Metric TSP](#)), approximation guarantee is possible
- however, even for Metric TSP, heuristics 'may' perform better

³Vazirani, V. Approximation algorithms. Springer Science & Business Media: 2013.

Local Search Methods

Local Search Methods

- computational problems arise due to size of the search space
- the idea of **local search methods** is to narrow down the space
- to this end, we define a **neighborhood** for variable assignments
- local search algorithms basically work as follows:
 - ① start with random assignment
 - ② replace current assignment with a 'reasonable' neighbor
 - ③ go back to 2 until no improvement is possible

Choosing a good Neighborhood

- use the whole search space as neighborhood
 - 1 the found solution will be optimal
 - 2 however, we do not gain anything computationally
- use empty set as neighborhood
 - 1 our algorithm will run in constant time
 - 2 however, the solution is just a random guess
- a good neighborhood balances between runtime and solution quality
- however, usually local search algorithms will find only locally optimal solutions (optimal in local neighborhood)

Swap-Neighborhood for TSP

- **Swap-Neighborhood:** select one variable and swap value with left neighbor
- each assignment has $n - 1 = \Theta(n)$ neighbors

Example: (2, 3, 4, 5) has neighbors

- (3, 2, 4, 5)
- (2, 4, 3, 5)
- (2, 3, 5, 4)
- (5, 3, 4, 2)

Transposition-Neighborhoods for TSP

- **Transposition-Neighborhood:** swap values of two arbitrary variables
- each assignment has $\binom{n-1}{2} = \frac{(n-1)(n-2)}{2} = \Theta(n^2)$ neighbors

Example: (2, 3, 4, 5) has neighbors

- (3, 2, 4, 5)
- (4, 3, 2, 5)
- (5, 3, 4, 2)
- (2, 4, 3, 5)
- (2, 5, 4, 3)
- (2, 3, 5, 4)

- the easiest local search method is **Hill-Climbing**
- it chooses the **steepest descent** neighbor (thus hill-climbing) until no improvement is possible anymore

Hill-Climbing

A \leftarrow random assignment

repeat

 if neighborhood of *A* contains better assignment *A'*

A \leftarrow best assignment in neighborhood

 else

 return *A*

Results for Random TSP with 15 nodes $\{0, 1, \dots, 14\}$

Initial assignment: (1,2,3,4,5,6,7,8,9,10,11,12,13,14)

Hill Climbing with Swap Neighborhood

Found local optimum after 5 iterations.

Cost: 52.0

Hill Climbing with Transposition Neighborhood

Found local optimum after 6 iterations.

Cost: 38.0

Initial assignment: (6,3,4,1,8,5,2,11,14,10,13,7,9,12)

Hill Climbing with Swap Neighborhood

Found local optimum after 6 iterations.

Cost: 47.0

Hill Climbing with Transposition Neighborhood

Found local optimum after 5 iterations.

Cost: 33.0

- Transposition Neighborhood gives us often better solutions than Swap Neighborhood
- this makes intuitively sense, since it allows to explore the search space better
- on the other hand, each iteration can take significantly more time (quadratic vs. linear)
- indeed, exploring the whole neighborhood can be expensive in large domains

First-Choice Hill-Climbing

- Hill-Climbing can take a lot of time if neighborhoods are large
- **First-Choice Hill-Climbing** stops iterating neighbors as soon as a better neighbor is found

First-Choice Hill-Climbing

$A \leftarrow \text{random assignment}$

repeat

$A^* \leftarrow A$

for neighbors A' of A^*

if A' improves A^*

$A^* \leftarrow A'$

skip loop

if A^* did not change

return A

Demo

Run	Algorithm	Neighborhood	Iterations	Cost	Time (ms)	Time/Iteration
1	HC	Swap	26	331	15	0.58
	HC	Transp	42	135	123	2.93
	FCHC	Swap	40	337	0	0
	FCHC	Transp	132	140	78	0.59
2	HC	Swap	35	328	0	0
	HC	Transp	35	159	38	1.01
	FCHC	Swap	48	337	15	0.31
	FCHC	Transp	156	145	69	0.44
3	HC	Swap	25	347	0	0
	HC	Transp	39	138	31	0.79
	FCHC	Swap	33	344	0	0
	FCHC	Transp	120	145	54	0.45

Results for Random TSP with 100 nodes, runs started from different initial assignments

- FCHC can sometimes find better local optima than HC
- the reason is that **locally suboptimal choices** can lead to better regions in search space
- two further improve **exploration of search space**, we can
 - ① allow suboptimal (or even worse) moves in search space
(e.g. *Stochastic Hill-Climbing*, *Simulated Annealing*)
 - ② consider multiple search paths in parallel
(e.g. *Parallel Hill-Climbing*, *Local Beam Search*)

- Local Search methods can be seen as metaheuristics
- choice of neighborhood balances between exploration (large) and efficiency (small)
- one of the fastest local search method is Hill Climbing
- in terms of solution quality, it is often outperformed by other metaheuristics that provide better exploration of search space
- however, Hill Climbing is often used in hybrid algorithms because of its efficiency (e.g. memetic algorithms)

Programming Task

Knapsack Problem



- **Problem:** given items I_1, \dots, I_n with weights w_1, \dots, w_n , values v_1, \dots, v_n and a weight limit W , maximize value
- **Examples**
 - invest budget in assets with maximum expected return
 - cut sheet metal into pieces of specified size, minimize waste
 - load truck with items of maximal value
 - subproblem (resource allocation) in many other problems

Knapsack Problem Formalized

- given n items l_1, \dots, l_n with
 - weights w_1, \dots, w_n
 - values v_1, \dots, v_n
 - and a weight limit W
- select items to meet weight limit and maximize value
- optimization problem
 - Variables: n variables with domain $\{0, 1\}$
($x_i = 1$ iff we select l_i)
 - Constraints: $w(x_1, \dots, x_n) = \sum_{i=1}^n w_i \cdot x_i \leq W$
 - Objective function: $f(x_1, \dots, x_n) = \sum_{i=1}^n v_i \cdot x_i$

Feasible Solutions

W = 100

w = 40, v=20

Pattern 1

w = 25, v=10

Pattern 2

w = 15, v=5

Pattern 3

	Pattern 1		Pattern 2				Pattern 3					
	x ₁	x ₂	x ₃	x ₄	x ₅	x ₆	x ₇	x ₈	x ₉	x ₁₀	x ₁₁	x ₁₂
w	40	40	25	25	25	25	15	15	15	15	15	15
v	20	20	10	10	10	10	5	5	5	5	5	5

- $v_1 = (1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0)$, $w(v_1) = 105 > 100$ (infeasible)
- $v_2 = (0, 0, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0)$, $w(v_2) = 100$, $f(v_2) = 40$ (feasible)
- $v_3 = (1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1)$, $w(v_3) = 95$, $f(v_3) = 45$ (feasible)

Knapsack Problem Programming Task

- define a 'small' and a 'large' neighborhood (don't forget to assure feasibility)
- implement Hill-Climbing and First-Choice Hill-Climbing
- make experiments with both implementations for some Knapsack instances and different initial assignments
- document your findings and prepare a small presentation (5-10 minutes)

Complexity of Knapsack Problem

- Knapsack Problem is NP-complete
- however, as opposed to TSP, it is only **weakly NP-hard** (there exists a 'pseudo-polynomial' algorithm)
- in particular, it **can be approximated efficiently** to an arbitrary degree (there exists an 'FPTAS')
- see (Vazirani 2013)⁴, Chapter 8 for more details

⁴Vazirani, V. Approximation algorithms. Springer Science & Business Media: 2013.

Questions

MY HOBBY: EMBEDDING NP-COMPLETE PROBLEMS IN RESTAURANT ORDERS

CHOTCHKIES RESTAURANT	
~ APPETIZERS ~	
MIXED FRUIT	2.15
FRENCH FRIES	2.75
SIDE SALAD	3.35
HOT WINGS	3.55
MOZZARELLA STICKS	4.20
SAMPLER PLATE	5.80
~ SANDWICHES ~	
BARBECUE	6.55



⁵ https://imgs.xkcd.com/comics/np_complete.png